

main.py

```
import time
import cv2
import numpy as np
import pyautogui
import pygetwindow as gw
import torch
from ai import GettingOverItAI
from constants import GAME_TITLE, BATCH_SIZE, SAVE_INTERVAL, BUFFER_MAXLEN
from models import SimpleModel, CuriosityModel
from preprocessing import preprocess_frame
from buffer import ReplayBuffer
from torch import optim

# Initialize the device
device = torch.device("cuda" if torch.cuda.is_available() else "cpu")

# Game details
window = gw.getWindowsWithTitle(GAME_TITLE)[0]
x, y, width, height = window.left, window.top, window.width, window.height

# The number of color channels in the input images
# This is 3 for RGB images or 1 for grayscale images
channels = 3 # Replace with the number of color channels in your preprocessed
images

# The number of frames stacked together
STACK_SIZE = 4 # This is a common value, but you should replace it with the actual
number you're using

input_shape = (STACK_SIZE * channels, height, width)
model = SimpleModel(input_shape)

def main():
    curiosity_model = CuriosityModel()
    ai = GettingOverItAI(model, curiosity_model)
    replay_buffer = ReplayBuffer(BUFFER_MAXLEN)

    last_save_time = time.time()

    # Define an optimizer
    optimizer = optim.Adam(model.parameters())

    while True:
        try:
            # Capture a screenshot of the game
            game_screen = pyautogui.screenshot(region=(x, y, width, height))
            game_screen_np = np.array(game_screen)
```

```

        # Preprocess the game screen
        preprocessed_screen_np = preprocess_frame(game_screen_np)

        # Convert the 12-channels image to multiple 3-channels (RGB) images for
visualization
        for i in range(0, preprocessed_screen_np.shape[2], 3): # Assume the
shape is (height, width, 12)
            visualization = preprocessed_screen_np[:, :, i:i + 3] # Take three
channels at a time

            # Scale the visualization to range 0-255 if it's not already in
that range
            visualization = cv2.normalize(visualization, None, alpha=0.0,
beta=255.0, norm_type=cv2.NORM_MINMAX,
                                         dtype=0)
            cv2.imshow(f'AI View Channels {i + 1}-{i + 3}', visualization)
            cv2.waitKey(1) # Refresh the display every 1 millisecond

        # Convert the numpy array to a tensor and move it to the device
        state_tensor =
torch.from_numpy(preprocessed_screen_np).float().to(device)

        # Convert the tensor back to a numpy array
        state = state_tensor.cpu().numpy()

        # Feed the preprocessed screen into the AI model
        predicted_action = ai.predict(state)

        # Save the current screen and action to the buffer
        replay_buffer.push(preprocessed_screen_np, predicted_action, 0, 0,
False)

        if len(replay_buffer) >= BATCH_SIZE:
            # When replay buffer has enough experiences
            state, action, reward, next_state, done =
replay_buffer.sample(BATCH_SIZE)

            # Set the gradients to zero
            optimizer.zero_grad()

            # Compute the loss
            loss = ai.compute_loss(state, action, reward, next_state, done)

            # Backpropagation
            loss.backward()

            # Optimization step
            optimizer.step()

```

```

        # Save the model at a regular interval
        if time.time() - last_save_time >= SAVE_INTERVAL:
            ai.save_model()
            last_save_time = time.time()

        time.sleep(0.1)

    except Exception as e:
        print(f"An error occurred: {e}")
        break

    finally:
        cv2.destroyAllWindows()

ai.save_model()

if __name__ == "__main__":
    main()

```

ai.py

```

import logging
import cv2
import numpy as np
import pyautogui
import torch
import torch.nn as nn
import torch.optim as optim
from skopt import gp_minimize
from torch import device
from torchvision.transforms import ToTensor
from buffer import ReplayBuffer
from constants import BUFFER_MAXLEN, EPSILON, EPSILON_DECAY, EPSILON_MIN,
OUTPUT_SIZE, BATCH_SIZE
from preprocessing import preprocess_frame, PLAYER_COLOR

logging.basicConfig(filename='ai.log', level=logging.ERROR)

# Initialize the device
device = torch.device("cuda" if torch.cuda.is_available() else "cpu")

class GettingOverItAI:
    def __init__(self, model, curiosity_model):
        self.model = model

```

```

self.curiosity_model = curiosity_model
self.loss_function = nn.MSELoss()
self.optimizer = optim.Adam(self.model.parameters(), lr=0.00025)
self.curiosity_optimizer = optim.Adam(self.curiosity_model.parameters(),
lr=0.00025)
self.transform = ToTensor()
self.total_reward = 0.0
self.previous_position = None
self.buffer = ReplayBuffer(BUFFER_MAXLEN)
self.epsilon = EPSILON
self.steps_since_progress = 0
self.max_position = None
self.previous_mouse_position = None
self.previous_position = 0.0
self.gamma = 0.99 # Discount factor for future rewards

```

```
@staticmethod
```

```
def get_screen_shot():
```

```

    try:
        # Capture the entire screen
        screen = np.array(pyautogui.screenshot())
        # Convert the image from BGR to RGB color space
        screen = cv2.cvtColor(screen, cv2.COLOR_BGR2RGB)
        return screen
    except Exception as e:
        logging.error(f'Error capturing screen: {e}')

```

```
@staticmethod
```

```
def perform_mouse_movement(action):
```

```

    try:
        # Get the current mouse position
        current_position = np.array(pyautogui.position())
        # Calculate the new mouse position
        new_position = current_position + action
        # Ensure the new position is within the screen bounds
        new_position = np.clip(new_position, 0, np.array(pyautogui.size()))
        # Move the mouse to the new position
        pyautogui.moveTo(new_position[0], new_position[1], duration=0.1)
    except Exception as e:
        logging.error(f'Error performing mouse movement: {e}')

```

```
def simulate_game(self, num_episodes=50):
```

```

    total_reward = 0.0
    self.previous_position = 0 # Initialize to a default value

```

```
    for _ in range(num_episodes):
```

```
        done = False
```

```

        # Get the initial game screen
        game_screen = self.get_screen_shot()

```

```

# Preprocess the game screen
state = preprocess_frame(game_screen)

while not done:
    action = self.get_mouse_movement()
    # Execute the action
    self.perform_mouse_movement(action)

    # Get the next game screen
    next_game_screen = self.get_screen_shot()

    # Preprocess the next game screen
    next_state = preprocess_frame(next_game_screen)

    # Determine if the game is "done"
    current_position = self.get_player_position(state)
    done = self.is_game_done(current_position)

    reward = self.get_reward(current_position)
    self.add_experience_to_buffer(state, action, reward, next_state,
done)

    self.train(BATCH_SIZE)

    # Update previous_position
    self.previous_position = current_position

    state = next_state
    total_reward += reward

average_reward = total_reward / num_episodes
return average_reward

def compute_loss(self, states, actions, rewards, next_states, dones):
    # Convert data to tensors
    states = torch.FloatTensor(states).to(device)
    next_states = torch.FloatTensor(next_states).to(device)
    actions = torch.LongTensor(actions).unsqueeze(-1).to(device) # Ensure that
actions are a column vector
    rewards = torch.FloatTensor(rewards).to(device)
    dones = torch.FloatTensor(dones).to(device)

    # Compute Q-values for the current states and next states
    curr_q_values = self.model(states)
    curr_q = curr_q_values.gather(1, actions).squeeze(-1) # Gather along the
action dimension
    next_q_values = self.model(next_states)
    next_q = next_q_values.max(1)[0]

    # Compute target Q-values

```

```

target_q = rewards + (1 - dones) * self.gamma * next_q

# Compute loss
loss = self.loss_function(curr_q, target_q.detach())
return loss

def objective(self, params):
    # Unpack the parameters
    num_episodes, = params

    # Simulate the game
    average_reward = self.simulate_game(num_episodes)

    # We're minimizing the objective function, so return the negative of the
reward    return -average_reward

def optimize(self):
    # Optimize the AI
    res = gp_minimize(self.objective, [(1, 100)], n_calls=50, random_state=0)

    # Print the result
    print(f'Best reward: {-res.fun}')

def get_mouse_movement(self):
    if np.random.rand() < self.epsilon:
        # Perform a random action
        action = np.random.randint(OUTPUT_SIZE)
    else:
        # Get the best action according to the model
        with torch.no_grad():
            state_tensor =
torch.from_numpy(self.previous_position).float().unsqueeze(0).to(device)
            q_values = self.model(state_tensor)
            action = torch.argmax(q_values).item()

    # Decay the epsilon value
    self.epsilon = max(self.epsilon * EPSILON_DECAY, EPSILON_MIN)
    return action

def get_player_position(self, state):
    # Function to determine the player's current position
    # This depends on the specifics of your game and may need to be adjusted
    player_pixels = np.where(np.all(state == PLAYER_COLOR / 255, axis=-1))
    if player_pixels[0].size > 0:
        return np.array([player_pixels[0].mean(), player_pixels[1].mean()])
    else:
        return self.previous_position

def is_game_done(self, current_position):

```

```

    # Function to determine if the game is "done" or not
    # This will need to be adjusted based on the specifics of your game
    done = False
    if self.previous_position is not None:
        # Consider the game done if the player has not made progress in the
last 200 steps
        self.steps_since_progress += 1
        if np.any(current_position > self.previous_position):
            self.steps_since_progress = 0
        if self.steps_since_progress > 200:
            done = True
    return done

def get_reward(self, current_position):
    # Function to calculate the reward based on the player's current position
    # This will need to be adjusted based on the specifics of your game
    if self.previous_position is not None:
        reward = np.maximum(current_position - self.previous_position,
0.0).sum()
    else:
        reward = 0.0
    return reward

def add_experience_to_buffer(self, state, action, reward, next_state, done):
    # Adds an experience to the replay buffer
    state = torch.from_numpy(state).float().to(device)
    next_state = torch.from_numpy(next_state).float().to(device)
    self.buffer.push(state, action, reward, next_state, done)

def predict(self, state):
    with torch.no_grad():
        state_tensor = torch.from_numpy(state).float().unsqueeze(0).to(device)
        q_values = self.model(state_tensor)
        action = torch.argmax(q_values).item()
    return action

def train(self, batch_size):
    # Trains the AI based on experiences from the replay buffer
    if len(self.buffer) >= batch_size:
        states, actions, rewards, next_states, dones =
self.buffer.sample(batch_size)
        targets = self.model(states).tolist()
        next_q_values = self.model(next_states).detach().numpy().max(axis=1)
        for i, done in enumerate(dones):
            if done:
                targets[i][actions[i]] = rewards[i]
            else:
                targets[i][actions[i]] = rewards[i] + next_q_values[i]
        targets = torch.tensor(targets).to(device)
        self.optimizer.zero_grad()

```

```

        loss = self.loss_function(self.model(states), targets)
        loss.backward()
        self.optimizer.step()

    def save_model(self, model_file_path='model.pth',
curiosity_model_file_path='curiosity_model.pth'):
        torch.save(self.model.state_dict(), model_file_path)
        torch.save(self.curiosity_model.state_dict(), curiosity_model_file_path)

```

```

models.py
import torch
from constants import OUTPUT_SIZE
import torch.nn as nn
import numpy as np

```

```

class SimpleModel(nn.Module):
    def __init__(self, input_shape): # Remove the default value
        super().__init__()
        self.conv = nn.Sequential(
            nn.Conv2d(input_shape[0], 32, kernel_size=5, stride=2),
nn.BatchNorm2d(32), nn.ReLU(),
            nn.Conv2d(32, 64, kernel_size=5, stride=2), nn.BatchNorm2d(64),
nn.ReLU(),
            nn.Conv2d(64, 64, kernel_size=5, stride=2), nn.BatchNorm2d(64),
nn.ReLU(),
        )

        # Create a dummy input with the correct input shape
        dummy_input = torch.zeros(1, *input_shape)

        # Pass the dummy input through the conv layers to get the output size
        conv_output = self.conv(dummy_input)

        # Calculate the number of features from the output size
        num_features = int(np.prod(conv_output.size()))

        self.fc = nn.Sequential(
            nn.Linear(num_features, 512), nn.ReLU(),
            nn.Linear(512, OUTPUT_SIZE)
        )

    def forward(self, x):
        x = self.conv(x)
        x = x.view(x.size(0), -1) # Flatten the tensor
        return self.fc(x)

```



```

class CuriosityModel(nn.Module):
    def __init__(self):
        super().__init__()
        self.inverse_model = nn.Sequential(
            nn.Linear(4480*2, 512), nn.ReLU(),
            nn.Linear(512, OUTPUT_SIZE)
        )
        self.forward_model = nn.Sequential(
            nn.Linear(4480+OUTPUT_SIZE, 512), nn.ReLU(),
            nn.Linear(512, 4480)
        )

    def forward(self, state, action, next_state):
        state_action = torch.cat([state, action], 1)
        predicted_next_state = self.forward_model(state_action)

        state_next_state = torch.cat([state, next_state], 1)
        predicted_action = self.inverse_model(state_next_state)

        return predicted_action, predicted_next_state

```

constants.py

```

# Game details
GAME_TITLE = 'Getting Over It'

# Hyperparameters
OUTPUT_SIZE = 2
BATCH_SIZE = 64
EPSILON = 0.9
EPSILON_DECAY = 0.999
EPSILON_MIN = 0.1
BUFFER_MAXLEN = int(1e5)
REWARD_THRESHOLD = 100
SAVE_INTERVAL = 600 # Save model every 10 minutes

# Miscellaneous
MODEL_SAVE_PATH = 'model_save.pth'

```

preprocessing.py

```

from collections import deque

```

```

import cv2
import numpy as np
import pyautogui
import pygetwindow as gw
import matplotlib.pyplot as plt

# Game details
game_title = 'Getting Over It'
window = gw.getWindowsWithTitle(game_title)[0]
x, y, width, height = window.left, window.top, window.width, window.height

# Frame stacking
stack_size = 4 # Adjust this based on your game's temporal dynamics
frame_stack = deque(maxlen=stack_size)

# Define the color for each object in the abstracted version
PLAYER_COLOR = np.array([255, 0, 0]) # Red
obstacle_color = np.array([0, 255, 0]) # Green
environment_color = np.array([0, 0, 255]) # Blue

# Data augmentation
noise_std = 0.1 # Adjust this based on the scale and nature of your data

def add_noise(frame):
    noise = np.random.normal(0, noise_std, frame.shape)
    frame += noise
    frame = np.clip(frame, 0, 1) # Ensure the added noise doesn't cause data to go
    out of range
    return frame

def rotate_frame(frame, angle):
    # Get image height and width
    (h, w) = frame.shape[:2]

    # Define the center of the image
    center = (w / 2, h / 2)

    # Perform the rotation
    m = cv2.getRotationMatrix2D(center, angle, 1.0)
    rotated = cv2.warpAffine(frame, m, (w, h))

    return rotated

def visualize_mask(mask, title):
    plt.figure(figsize=(10, 10))
    plt.imshow(mask, cmap='hot')
    plt.title(title)

```

```
plt.show()
```

```
def visualize_histogram(data, title):  
    plt.figure(figsize=(10, 10))  
    plt.hist(data.ravel(), bins=256, color='orange', )  
    plt.title(title)  
    plt.show()
```

```
def visualize_frame_stack(frame_stack):  
    for i, frame in enumerate(frame_stack):  
        # Transpose the frame back to its original shape  
        frame = np.transpose(frame, (1, 2, 0))  
  
        # Rescale the frame to the range 0-255  
        frame = (frame * 255).astype(np.uint8)  
  
        plt.figure(figsize=(10, 10))  
        plt.imshow(frame, cmap='gray')  
        plt.title(f'Frame {i+1}')  
        plt.show()
```

```
# Define the size for the low-res, pixelated version  
low_res_size = (64, 64) # Adjust this based on your needs
```

```
def preprocess_frame(frame, frame_stack):  
    # Convert the frame to HSV color space  
    hsv = cv2.cvtColor(frame, cv2.COLOR_BGR2HSV)  
  
    # Define the color range for each object  
    player_range = (np.array([0, 100, 100]), np.array([10, 255, 255])) # Red range  
    obstacle_range = (np.array([35, 100, 100]), np.array([85, 255, 255])) # Green  
range  
    environment_range = (np.array([110, 100, 100]), np.array([130, 255, 255])) #  
Blue range  
  
    # Create a binary mask for each object  
    player_mask = cv2.inRange(hsv, player_range[0], player_range[1])  
    obstacle_mask = cv2.inRange(hsv, obstacle_range[0], obstacle_range[1])  
    environment_mask = cv2.inRange(hsv, environment_range[0], environment_range[1])  
  
    # Combine the masks into a single image  
    preprocessed = np.zeros_like(frame)  
    player_color = (255, 255, 255) # example RGB color for the player  
    preprocessed[player_mask > 0] = player_color  
    preprocessed[obstacle_mask > 0] = obstacle_color  
    preprocessed[environment_mask > 0] = environment_color
```

```

# Resize the preprocessed frame to a lower resolution
preprocessed = cv2.resize(preprocessed, low_res_size)

# Normalize the preprocessed frame
preprocessed = preprocessed / 255.0

# Add noise for data augmentation
preprocessed = add_noise(preprocessed)

# Transpose the frame to have channels as the first dimension
preprocessed = np.transpose(preprocessed, (2, 0, 1))

# Add the preprocessed frame to the frame stack
frame_stack.append(preprocessed)

# Stack the frames along the channel dimension
if len(frame_stack) < stack_size:
    # If stack is not full, pad with zeros
    stacked_frames = np.concatenate(
        [np.zeros(((stack_size - len(frame_stack)) * preprocessed.shape[0],
                    preprocessed.shape[1], preprocessed.shape[2])),
        *frame_stack], axis=0)
    else:
        # If stack is full or overfull, drop the oldest frames and use the newest
        ones
        stacked_frames = np.concatenate(list(frame_stack), axis=0)[-stack_size *
preprocessed.shape[0]:]

    return stacked_frames, frame_stack

def display_preprocessed_frame(frame):
    # Convert frame to a NumPy array if it's not already
    if not isinstance(frame, np.ndarray):
        frame = np.array(frame)

    # Undo the normalization and transpose back to the original shape
    display_frame = (frame * 255.0).astype(np.uint8)
    display_frame = np.transpose(display_frame, (1, 2, 0))

    # Use only the first three channels of the image
    if display_frame.shape[2] > 3:
        display_frame = display_frame[:, :, :3]

    # Create a figure with subplots
    fig, ax = plt.subplots(figsize=(10, 10))

    # Show the frame with a color map
    im = ax.imshow(display_frame, cmap='hot')

```

```

# Add a color bar
cbar = fig.colorbar(im, ax=ax)
cbar.set_label('Pixel Intensity')

# Add a title
ax.set_title('Preprocessed Frame')

# Display the figure
plt.show()

# Capture a screenshot of the game
while True:
    game_screen = pyautogui.screenshot(region=(x, y, width, height))
    game_screen_np = np.array(game_screen)
    preprocessed_frame, frame_stack = preprocess_frame(game_screen_np, frame_stack)

    display_preprocessed_frame(preprocessed_frame) # Show the processed frame

    # Break the loop when 'q' is pressed
    if cv2.waitKey(1) & 0xFF == ord('q'):
        break

cv2.destroyAllWindows()

```

buffer.py

```

import random
from collections import deque
import numpy as np

class ReplayBuffer:
    def __init__(self, capacity):
        self.capacity = capacity
        self.buffer = deque(maxlen=capacity)

    def push(self, state, action, reward, next_state, done):
        state = np.expand_dims(state, 0)
        next_state = np.expand_dims(next_state, 0)

        self.buffer.append((state, action, reward, next_state, done))

    def sample(self, batch_size):
        state, action, reward, next_state, done = zip(*random.sample(self.buffer,

```

```
batch_size))
    return np.concatenate(state), action, reward, np.concatenate(next_state),
done

def __len__(self):
    return len(self.buffer)
```